

Environmental Health Big Data Analysis – R4ds (2) Data Wrangling and Programming

Il-Youp Kwak
Chung-Ang University



What we will learn

1. Data transformation
2. Data visualization
- 3. Data wrangling**
- 4. Functional programming with R**
- 5. Something useful**




Data Wrangling - tidyr

“Tidy datasets are all alike, but every messy dataset is messy in its own way.” — Hadley Wickham

Tidy data:

- **Every column is variable**
- **Every row is an observation.**
- **Every cell is a single value.**



Why tidyr?

- **If you have a consistent data structure, it's easier to learn the tools that work with it.**
- **Most built-in R functions work with vectors of values. That makes transforming tidy data feel particularly natural.**
- **dplyr, ggplot2, and all the other packages in the tidyverse are designed to work with tidy data.**



Example Data

Representation of the same data in multiple ways

table1

```
#> # A tibble: 6 x 4
#>   country      year cases population
#>   <chr>      <int> <int>      <int>
#> 1 Afghanistan 1999     745 19987071
#> 2 Afghanistan 2000    2666 20595360
#> 3 Brazil      1999   37737 172006362
#> 4 Brazil      2000   80488 174504898
#> 5 China       1999  212258 1272915272
#> 6 China       2000  213766 1280428583
```

table2

```
#> # A tibble: 12 x 4
#>   country      year type      count
#>   <chr>      <int> <chr>      <int>
#> 1 Afghanistan 1999 cases         745
#> 2 Afghanistan 1999 population 19987071
#> 3 Afghanistan 2000 cases         2666
#> 4 Afghanistan 2000 population 20595360
#> 5 Brazil      1999 cases         37737
#> 6 Brazil      1999 population 172006362
#> # ... with 6 more rows
```

table3

```
#> # A tibble: 6 x 3
#>   country      year rate
#>   * <chr>      <int> <chr>
#> 1 Afghanistan 1999 745/19987071
#> 2 Afghanistan 2000 2666/20595360
#> 3 Brazil      1999 37737/172006362
#> 4 Brazil      2000 80488/174504898
#> 5 China       1999 212258/1272915272
#> 6 China       2000 213766/1280428583
```

table4a # cases

```
#> # A tibble: 3 x 3
#>   country `1999` `2000`
#>   * <chr>      <int> <int>
#> 1 Afghanistan     745     2666
#> 2 Brazil          37737    80488
#> 3 China          212258    213766
```

table4b # population

```
#> # A tibble: 3 x 3
#>   country `1999` `2000`
#>   * <chr>      <int> <int>
#> 1 Afghanistan 19987071 20595360
#> 2 Brazil      172006362 174504898
#> 3 China      1272915272 1280428583
```



Pivoting

- **One variable might be spread across multiple columns. (Use `Pivot_longer()`)**

```
table4a # cases
#> # A tibble: 3 x 3
#>   country   `1999` `2000`
#> * <chr>     <int> <int>
#> 1 Afghanistan    745   2666
#> 2 Brazil       37737  80488
#> 3 China        212258 213766
```

```
table4b # population
#> # A tibble: 3 x 3
#>   country   `1999` `2000`
#> * <chr>     <int> <int>
#> 1 Afghanistan 19987071 20595360
#> 2 Brazil     172006362 174504898
#> 3 China     1272915272 1280428583
```

- **One observation might be scattered across multiple rows. (Use `Pivot_wider()`)**



Apply Pivot_longer() to table4a

```
table4a %>%  
  pivot_longer(c(`1999`, `2000`), names_to = "year", values_to = "cases")  
#> # A tibble: 6 x 3  
#>   country   year cases  
#>   <chr>    <chr> <int>  
#> 1 Afghanistan 1999    745  
#> 2 Afghanistan 2000   2666  
#> 3 Brazil      1999  37737  
#> 4 Brazil      2000  80488  
#> 5 China       1999 212258  
#> 6 China       2000 213766
```

country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

country	1999	2000
Afghanistan	745	2666
Brazil	37737	80488
China	212258	213766

table4



Apply Pivot_longer() and combine data

```
tidy4a <- table4a %>%  
  pivot_longer(c(`1999`, `2000`), names_to = "year", values_to = "cases")  
tidy4b <- table4b %>%  
  pivot_longer(c(`1999`, `2000`), names_to = "year", values_to = "population")  
left_join(tidy4a, tidy4b)  
#> Joining, by = c("country", "year")  
#> # A tibble: 6 x 4  
#>   country   year cases population  
#>   <chr>    <chr> <int>    <int>  
#> 1 Afghanistan 1999    745  19987071  
#> 2 Afghanistan 2000   2666  20595360  
#> 3 Brazil      1999  37737  172006362  
#> 4 Brazil      2000  80488  174504898  
#> 5 China       1999 212258 1272915272  
#> 6 China       2000 213766 1280428583
```

```
table4a # cases  
#> # A tibble: 3 x 3  
#>   country `1999` `2000`  
#> * <chr>    <int> <int>  
#> 1 Afghanistan    745   2666  
#> 2 Brazil      37737  80488  
#> 3 China      212258 213766
```

```
table4b # population  
#> # A tibble: 3 x 3  
#>   country `1999` `2000`  
#> * <chr>    <int> <int>  
#> 1 Afghanistan 19987071 20595360  
#> 2 Brazil      172006362 174504898  
#> 3 China      1272915272 1280428583
```



+ Code + Text

✓ RAM  Disk  Editing 

▼ R for Data Science 실습 2 - Data wrangling and programing

Data Wrangling (tidyr)

```
✓ [1] library(tidyverse)
1s

Warning message in system("timedatectl", intern = TRUE):
"running command 'timedatectl' had status 1"
── Attaching packages ─────────────────── tidyverse 1.3.1 ───

✓ ggplot2 3.3.5    ✓ purrr   0.3.4
✓ tibble  3.1.4    ✓ dplyr   1.0.7
✓ tidyr   1.1.3    ✓ stringr 1.4.0
✓ readr   2.0.1    ✓ forcats 0.5.1

── Conflicts ─────────────────────────── tidyverse_conflicts() ───
✗ dplyr::filter() masks stats::filter()
✗ dplyr::lag()     masks stats::lag()
```

↑ ↓ ↺ ✎ 📄 🗑️ ⋮

▼ Example datasets

```
[ ] table1

A tibble: 6 × 4
```

✓ 1s completed at 8:59 PM



Pivoting

- One variable might be spread across multiple columns. (Use `Pivot_longer()`)
- **One observation might be scattered across multiple rows. (Use `Pivot_wider()`)**

```
table2
#> # A tibble: 12 x 4
#>   country    year type      count
#>   <chr>    <int> <chr>    <int>
#> 1 Afghanistan 1999 cases      745
#> 2 Afghanistan 1999 population 19987071
#> 3 Afghanistan 2000 cases      2666
#> 4 Afghanistan 2000 population 20595360
#> 5 Brazil      1999 cases      37737
#> 6 Brazil      1999 population 172006362
#> # ... with 6 more rows
```



Apply Pivot_wider()

```
table2 %>%  
  pivot_wider(names_from = type, values_from = count)  
#> # A tibble: 6 x 4  
#>   country    year cases population  
#>   <chr>    <int> <int>    <int>  
#> 1 Afghanistan 1999    745  19987071  
#> 2 Afghanistan 2000   2666  20595360  
#> 3 Brazil      1999   37737 172006362  
#> 4 Brazil      2000   80488 174504898  
#> 5 China       1999  212258 1272915272  
#> 6 China       2000  213766 1280428583
```

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

table2

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583



0s

A tibble: 6 × 4

country	year	cases	population
<chr>	<chr>	<int>	<int>
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

Q: What is right_join? try ?right_join and study join methods

▼ Pivot_wider()

One observation might be scattered across multiple rows

```
[ ] table2
```

A tibble: 12 × 4

country	year	type	count
---------	------	------	-------

Separating and uniting

- **One column contains two variables (Use separate())**


```
table3
#> # A tibble: 6 x 3
#>   country    year rate
#> * <chr>    <int> <chr>
#> 1 Afghanistan 1999 745/19987071
#> 2 Afghanistan 2000 2666/20595360
#> 3 Brazil      1999 37737/172006362
#> 4 Brazil      2000 80488/174504898
#> 5 China       1999 212258/1272915272
#> 6 China       2000 213766/1280428583
```

- **Single variable is spread across multiple columns.
(Use unite())**



Apply separate()

```
table3 %>%  
  separate(rate, into = c("cases", "population"))  
#> # A tibble: 6 x 4  
#>   country   year cases population  
#>   <chr>    <int> <chr>   <chr>  
#> 1 Afghanistan 1999 745    19987071  
#> 2 Afghanistan 2000 2666   20595360  
#> 3 Brazil      1999 37737  172006362  
#> 4 Brazil      2000 80488  174504898  
#> 5 China       1999 212258 1272915272  
#> 6 China       2000 213766 1280428583
```



Separating and uniting


- One column contains two variables (Use `separate()`)
- **Single variable is spread across multiple columns.**
(Use `unite()`)

	country	century	year	rate
	<chr>	<chr>	<chr>	<chr>
1	Afghanistan	19	99	745/19987071
2	Afghanistan	20	00	2666/20595360
3	Brazil	19	99	37737/172006362
4	Brazil	20	00	80488/174504898
5	China	19	99	212258/1272915272
6	China	20	00	213766/1280428583



Apply unite()

```
table5 %>%  
  unite(new, century, year, sep = "")  
#> # A tibble: 6 x 3  
#>   country    new    rate  
#>   <chr>    <chr> <chr>  
#> 1 Afghanistan 1999 745/19987071  
#> 2 Afghanistan 2000 2666/20595360  
#> 3 Brazil      1999 37737/172006362  
#> 4 Brazil      2000 80488/174504898  
#> 5 China       1999 212258/1272915272  
#> 6 China       2000 213766/1280428583
```



country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	century	year	rate
Afghanistan	19	99	745 / 19987071
Afghanistan	20	00	2666 / 20595360
Brazil	19	99	37737 / 172006362
Brazil	20	00	80488 / 174504898
China	19	99	212258 / 1272915272
China	20	00	213766 / 1280428583

table6



+ Code + Text

✓ RAM
Disk Editing

Separating and Uniting

▼ separate()

One column contains two variables

```
[ ] table3
```

A tibble: 6 × 3

	country	year	rate
	<chr>	<int>	<chr>
1	Afghanistan	1999	745/19987071
2	Afghanistan	2000	2666/20595360
3	Brazil	1999	37737/172006362
4	Brazil	2000	80488/174504898
5	China	1999	212258/1272915272
6	China	2000	213766/1280428583

```
[ ] table3 %>% separate(rate, into = c("cases", "population"))
```

Functional Programming

- **For loops are quite verbose, and require quite a bit of bookkeeping code that is duplicated for every for loop.**
- **Functional programming offers tools to extract out this duplicated code, so each common for loop pattern gets its own function.**



For loops

Random data

```
df <- tibble(  
  a = rnorm(10),  
  b = rnorm(10),  
  c = rnorm(10),  
  d = rnorm(10)  
)
```

With for loops

```
output <- vector("double", ncol(df)) # 1. output  
for (i in seq_along(df)) {          # 2. sequence  
  output[[i]] <- median(df[[i]])    # 3. body  
}  
output  
#> [1] -0.24576245 -0.28730721 -0.05669771 0.14426335
```

To calculate median

```
median(df$a)  
#> [1] -0.2457625  
median(df$b)  
#> [1] -0.2873072  
median(df$c)  
#> [1] -0.05669771  
median(df$d)  
#> [1] 0.1442633
```



Three components

The **output**: `output <- vector("double", length(x))`.

Allocate sufficient space for the output.

(If you grow the for loop at each iteration using `c()` (for example), your for loop will be very slow)

The **sequence**: `i in seq_along(df)`.

what to loop over:

each run of the for loop will assign `i` to a different value from `seq_along(df)`

The **body**: `output[[i]] <- median(df[[i]])`.

This is the code that does the work.

```
output <- vector("double", ncol(df)) # 1. output
for (i in seq_along(df)) {          # 2. sequence
  output[[i]] <- median(df[[i]])    # 3. body
}
output
#> [1] -0.24576245 -0.28730721 -0.05669771 0.14426335
```



For loops vs. functionals

Possible to wrap up for loops in a function

```
col_mean <- function(df) {  
  output <- vector("double", ncol(df))  
  for (i in seq_along(df)) {  
    output[[i]] <- mean(df[[i]])  
  }  
  output  
}
```

```
col_median <- function(df) {  
  output <- vector("double", ncol(df))  
  for (i in seq_along(df)) {  
    output[[i]] <- median(df[[i]])  
  }  
  output  
}
```

```
col_sd <- function(df) {  
  output <- vector("double", ncol(df))  
  for (i in seq_along(df)) {  
    output[[i]] <- sd(df[[i]])  
  }  
  output  
}
```

```
col_summary <- function(df, fun) {  
  out <- vector("double", length(df))  
  for (i in seq_along(df)) {  
    out[i] <- fun(df[[i]])  
  }  
  out  
}  
col_summary(df, median)  
#> [1] -0.51850298 0.02779864 0.17295591 -0.61163819  
col_summary(df, mean)  
#> [1] -0.3260369 0.1356639 0.4291403 -0.2498034
```





+ Code + Text

✓ RAM  Disk  Editing 

✓ [25] 0s

Brazil	1999	37737/172006362
Brazil	2000	80488/174504898
China	1999	212258/1272915272
China	2000	213766/1280428583



Functional Programming

```
df <- tibble(  
  a = rnorm(10),  
  b = rnorm(10),  
  c = rnorm(10),  
  d = rnorm(10)  
)
```

```
[ ] median(df$a)  
#> [1] -0.2457625  
median(df$b)  
#> [1] -0.2873072  
median(df$c)  
#> [1] -0.05669771  
median(df$d)  
#> [1] 0.1442633
```

```
0.423909950150083  
0.0269074082896925  
-0.0806930728782541
```



The map function (purrr)

the purrr package provides a family of functions for looping patterns over a vector

- **map()** makes a list
- **map_lgl()** makes a logical vector
- **map_int()** makes an integer vector
- **map_dbl()** makes a double vector
- **map_chr()** makes a character vector

Alternatives: apply, lapply, etc



The map function (purrr)

```
df %>% map_dbl(mean)
#>      a      b      c      d
#> -0.3260369 0.1356639 0.4291403 -0.2498034
df %>% map_dbl(median)
#>      a      b      c      d
#> -0.51850298 0.02779864 0.17295591 -0.61163819
df %>% map_dbl(sd)
#>      a      b      c      d
#> 0.9214834 0.4848945 0.9816016 1.1563324
```

You can define a function in a map function

```
models <- mtcars %>%
  split(.$cyl) %>%
  map(function(df) lm(mpg ~ wt, data = df))
```

```
models <- mtcars %>%
  split(.$cyl) %>%
  map(~lm(mpg ~ wt, data = .))
```



To extract a component

```
models %>%  
  map(summary) %>%  
  map_dbl(~.$r.squared)  
#>      4      6      8  
#> 0.5086326 0.4645102 0.4229655
```

```
models %>%  
  map(summary) %>%  
  map_dbl("r.squared")  
#>      4      6      8  
#> 0.5086326 0.4645102 0.4229655
```

You can also use an integer to select elements by position

```
x <- list(list(1, 2, 3), list(4, 5, 6), list(7, 8, 9))  
x %>% map_dbl(2)  
#> [1] 2 5 8
```



+ Code + Text

✓ RAM  Disk  Editing 

✓ [43]
0s
0.81344836910122 · 0.405377190952006 · 0.520463400489272 · -0.242193831844869

▼ The map function (purrr)

the purrr package provides a family of functions for looping patterns over a vector
remind apply()

```
[ ] str(df)

tibble [10 × 4] (S3: tbl_df/tbl/data.frame)
 $ a: num [1:10] -0.97 0.586 -2.17 0.906 -0.997 ...
 $ b: num [1:10] 1.064 -1.168 -0.623 0.239 -1.17 ...
 $ c: num [1:10] -0.0722 -0.0892 0.4624 0.7497 0.983 ...
 $ d: num [1:10] 0.516 -0.508 0.301 -0.813 0.284 ...

[ ] df %>% map_dbl(mean)

a:      -0.0846391036597818 b:      -0.00240030324052032 c:      -0.0309176854543395 d:      -0.0832475108396539

[ ] df %>% map_dbl(median)

a:      0.423909950150083 b:      0.0269074082896925 c:      -0.0806930728782541 d:      0.0468444173239768

[ ] mtcars %>%
  split(.$cyl)
```

Something useful

- R markdown
- **Boxplot vs Violinplot**
- **Do not use pie chart?**
- **Effective 2D visualization with t-SNE and UMAP**



Table of contents

Welcome

1 Introduction

Explore

2 Introduction

3 Data visualisation

4 Workflow: basics

5 Data transformation

6 Workflow: scripts

7 Exploratory Data Analysis

8 Workflow: projects

Wrangle

9 Introduction

10 Tibbles

11 Data import

12 Tidy data

13 Relational data

14 Strings

15 Factors

16 Dates and times

27 R Markdown

27.1 Introduction

R Markdown provides an unified authoring framework for data science, combining your code, its results, and your prose commentary. R Markdown documents are fully reproducible and support dozens of output formats, like PDFs, Word files, slideshows, and more.

R Markdown files are designed to be used in three ways:

1. For communicating to decision makers, who want to focus on the conclusions, not the code behind the analysis.
2. For collaborating with other data scientists (including future you!), who are interested in both your conclusions, and how you reached them (i.e. the code).
3. As an environment in which to *do* data science, as a modern day lab notebook where you can capture not only what you did, but also what you were thinking.

R Markdown integrates a number of R packages and external tools. This means that help is, by-and-large, not available through `?` . Instead, as you work through this chapter, and use R Markdown in the future, keep these resources close to hand:

- R Markdown Cheat Sheet: *Help > Cheatsheets > R Markdown Cheat Sheet*,

On this page

27 R Markdown

27.1 Introduction

27.2 R Markdown basics

27.3 Text formatting with Markdown

27.4 Code chunks

27.5 Troubleshooting


27.6 YAML header

27.7 Learning more



View source 

Edit this page 

Violinplot shows the data density



Example





+ Code + Text




✓ RAM Disk | Editing

0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
0.24	Very Good	J	VVS2	62.8	57	336	3.94	3.96	2.48
0.24	Very Good	I	VVS1	62.3	57	336	3.95	3.98	2.47
0.26	Very Good	H	SI1	61.9	55	337	4.07	4.11	2.53
0.22	Fair	E	VS2	65.1	61	337	3.87	3.78	2.49
0.23	Very Good	H	VS1	59.4	61	338	4.00	4.05	2.39
0.30	Good	J	SI1	64.0	55	339	4.25	4.28	2.73
0.23	Ideal	J	VS1	62.8	56	340	3.93	3.90	2.46
0.22	Premium	F	SI1	60.4	61	342	3.88	3.84	2.33
0.31	Ideal	J	SI2	62.2	54	344	4.35	4.37	2.71
0.20	Premium	E	SI2	60.2	62	345	3.79	3.75	2.27
0.32	Premium	E	I1	60.9	58	345	4.38	4.42	2.68
0.30	Ideal	I	SI2	62.0	54	348	4.31	4.34	2.68
0.30	Good	J	SI1	63.4	54	351	4.23	4.29	2.70
0.30	Good	J	SI1	63.8	56	351	4.23	4.26	2.71
0.30	Very Good	J	SI1	62.7	59	351	4.21	4.27	2.66
0.30	Good	I	SI2	63.3	56	351	4.26	4.30	2.71
0.22	Very Good	E	VS2	62.8	55	352	3.95	3.92	2.48




✓ 1s completed at 10:29 PM





Do not use pie chart?




Same Information from Bar Plot



Human Eyes are not Good at Evaluating Sizes

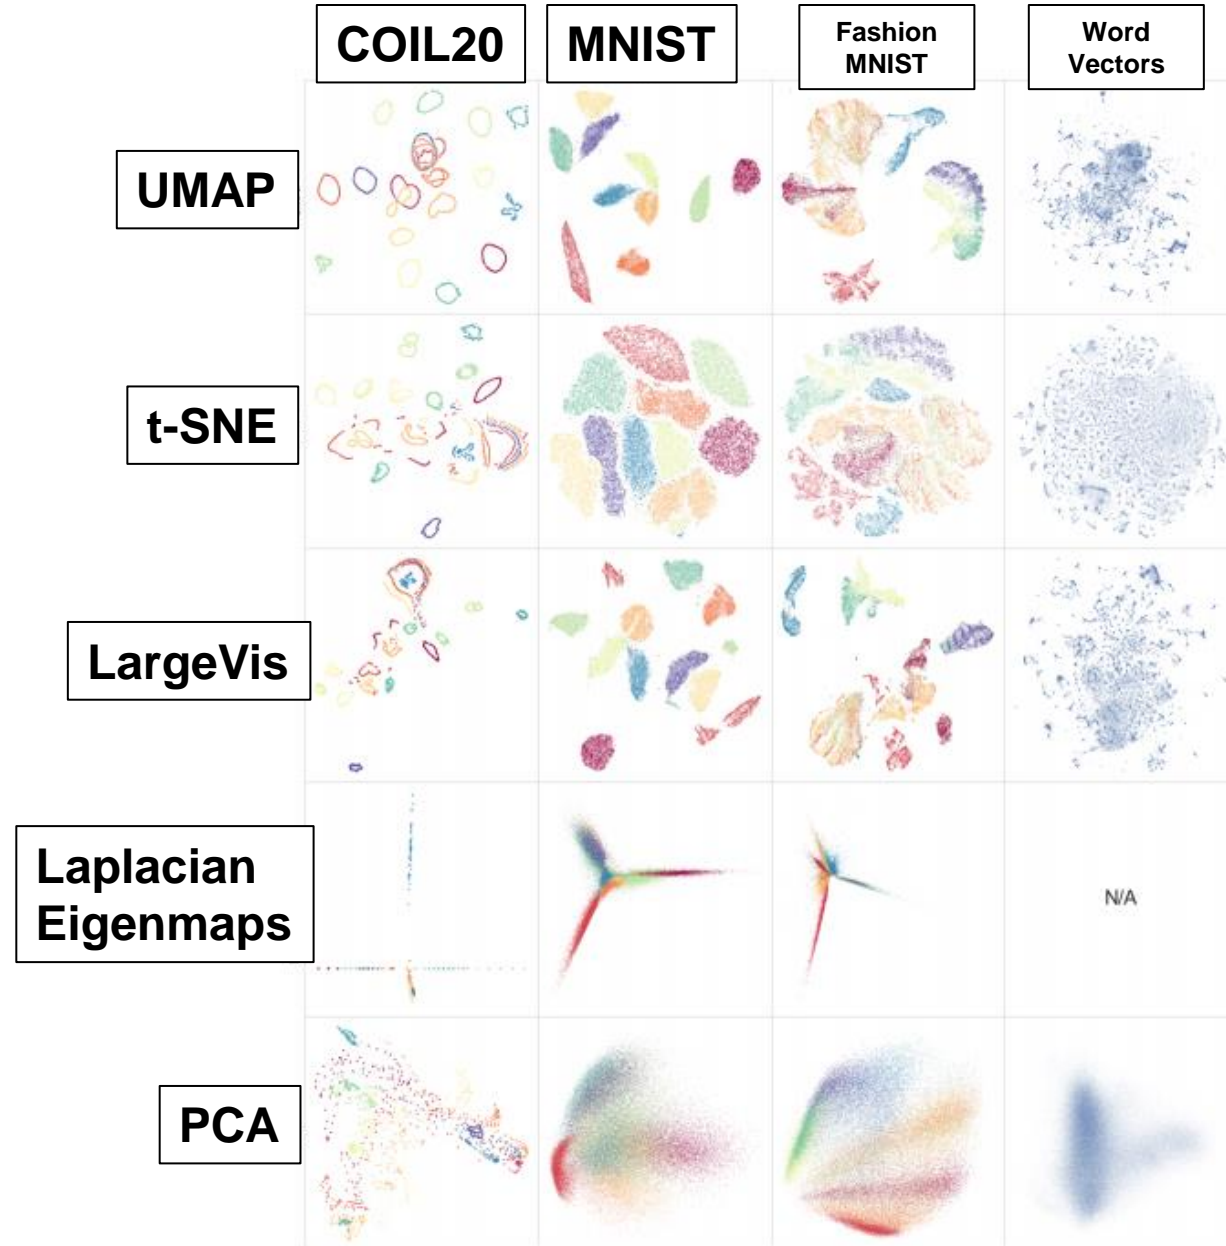


3d Plot even Worse

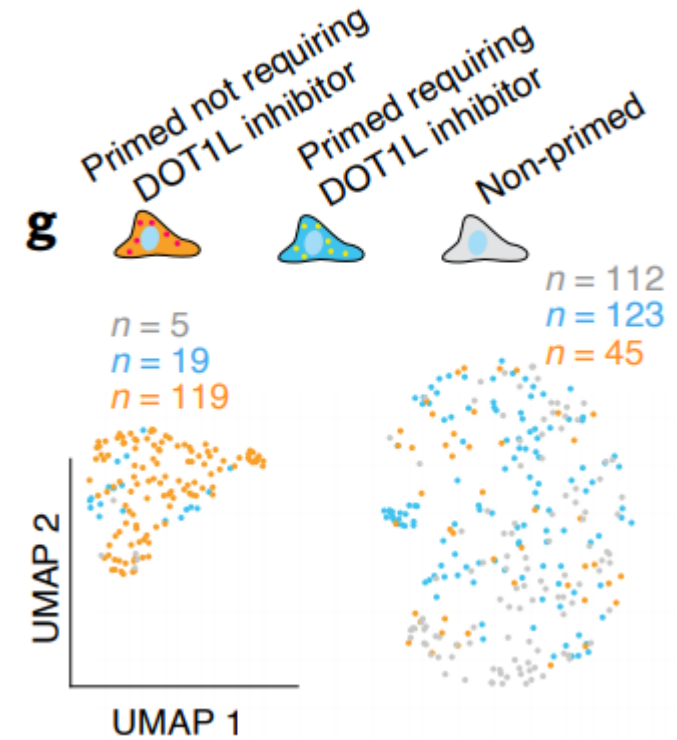


Effective 2D visualization with t-SNE and UMAP

When visualizing multidimensional data in 2D, t-SNE and UMAP works well



Ex.






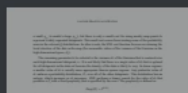
1



2



3



Visualizing Data using t-SNE

Laurens van der Maaten

TiCC

Tilburg University

P.O. Box 90153, 5000 LE Tilburg, The Netherlands

LVDMAATEN@GMAIL.COM

Geoffrey Hinton

Department of Computer Science

University of Toronto

6 King's College Road, M5S 3G4 Toronto, ON, Canada

HINTON@CS.TORONTO.EDU

Editor: Yoshua Bengio

Abstract

We present a new technique called “t-SNE” that visualizes high-dimensional data by giving each datapoint a location in a two or three-dimensional map. The technique is a variation of Stochastic Neighbor Embedding (Hinton and Roweis, 2002) that is much easier to optimize, and produces significantly better visualizations by reducing the tendency to crowd points together in the center of the map. t-SNE is better than existing techniques at creating a single map that reveals structure at many different scales. This is particularly important for high-dimensional data that lie on several different, but related, low-dimensional manifolds, such as images of objects from multiple classes seen from multiple viewpoints. For visualizing the structure of very large data sets, we show how t-SNE can use random walks on neighborhood graphs to allow the implicit structure of all of the data to influence the way in which a subset of the data is displayed. We illustrate the performance of t-SNE on a wide variety of data sets and compare it with many other non-parametric visualization techniques, including Sammon mapping, Isomap, and Locally Linear Embedding. The visualizations produced by t-SNE are significantly better than those produced by the other techniques on almost all of the data sets.

Thank you! 😊

